

predictor_lr

March 12, 2024

```
[14]: import numpy as np
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error
import matplotlib.pyplot as plt
import seaborn as sns

[15]: def load_dataset(file_path):
    """Load the dataset from a CSV file"""
    return pd.read_csv(file_path)

[16]: def preprocess_data(data):
    """Preprocess the dataset by handling missing values"""
    # Fill missing values with the mean of the respective column
    data = data.fillna(data.mean())
    return data

[17]: def explore_dataset(data):
    """Display the first few rows and information of the dataset"""
    print("First few rows of the dataset:")
    print(data.head())
    print("\nDataset information:")
    print(data.info())

[18]: def visualize_data(data):
    """Visualize the data"""
    # Pairplot to visualize relationships between features
    sns.pairplot(data, x_vars=data.columns[:-1], y_vars=['medv'],
    ↪kind='scatter')
    plt.title("Pairplot of Features vs. Target")
    plt.show()

    # Heatmap to visualize correlations between features
    plt.figure(figsize=(12, 8))
    sns.heatmap(data.corr(), annot=True, cmap='coolwarm', linewidths=0.5)
    plt.title("Correlation Heatmap")
    plt.show()
```

```
[19]: def split_data(data):
        """Split the dataset into features (X) and target (y)"""
        X = data.drop('medv', axis=1)
        y = data['medv']
        return X, y

[20]: def train_model(X_train, y_train):
        """Train a linear regression model"""
        model = LinearRegression()
        model.fit(X_train, y_train)
        return model

[21]: def evaluate_model(model, X_test, y_test):
        """Evaluate the trained model"""
        y_pred = model.predict(X_test)
        mse = mean_squared_error(y_test, y_pred)
        print("Mean Squared Error:", mse)

[22]: def visualize_predictions(y_test, y_pred):
        """Visualize the actual vs. predicted prices"""
        plt.figure(figsize=(10, 6))
        plt.scatter(y_test, y_pred, color='blue')
        plt.plot([y_test.min(), y_test.max()], [y_test.min(), y_test.max()],
        ↪linestyle='--', color='red')
        plt.xlabel("Actual Price")
        plt.ylabel("Predicted Price")
        plt.title("Actual vs. Predicted Prices")
        plt.show()

[23]: def predict_new_house_price(model, new_house_features):
        """Predict the price of a new house"""
        predicted_price = model.predict(new_house_features)
        formatted_price = "${:,.2f}".format(predicted_price[0])
        print("Predicted Price of the House:", formatted_price)

[24]: def main():
        # Load the dataset
        file_path = 'dataset/BostonHousing.csv'
        boston = load_dataset(file_path)

        # Preprocess the dataset
        boston = preprocess_data(boston)

        # Explore the dataset
        explore_dataset(boston)
        visualize_data(boston)
```

```

# Split the data into training and testing sets
X, y = split_data(boston)
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
↳random_state=42)

# Train the linear regression model
model = train_model(X_train, y_train)

# Evaluate the model
evaluate_model(model, X_test, y_test)

# Make predictions
y_pred = model.predict(X_test)
visualize_predictions(y_test, y_pred)

# Predict the price of a new house
new_house_features = np.array([[0.00632, 18, 2.31, 0, 0.538, 6.575, 65.2, 4.
↳09, 1, 296, 15.3, 396.9, 4.98]])
predict_new_house_price(model, new_house_features)

```

```

[25]: if __name__ == "__main__":
# Hide warning concerning feature names
import warnings
warnings.filterwarnings("ignore", message="X does not have valid feature_
↳names, but LinearRegression was fitted with feature names")

main()

```

First few rows of the dataset:

	crim	zn	indus	chas	nox	rm	age	dis	rad	tax	ptratio	\
0	0.00632	18.0	2.31	0	0.538	6.575	65.2	4.0900	1	296	15.3	
1	0.02731	0.0	7.07	0	0.469	6.421	78.9	4.9671	2	242	17.8	
2	0.02729	0.0	7.07	0	0.469	7.185	61.1	4.9671	2	242	17.8	
3	0.03237	0.0	2.18	0	0.458	6.998	45.8	6.0622	3	222	18.7	
4	0.06905	0.0	2.18	0	0.458	7.147	54.2	6.0622	3	222	18.7	

	b	lstat	medv
0	396.90	4.98	24.0
1	396.90	9.14	21.6
2	392.83	4.03	34.7
3	394.63	2.94	33.4
4	396.90	5.33	36.2

Dataset information:

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 506 entries, 0 to 505
Data columns (total 14 columns):

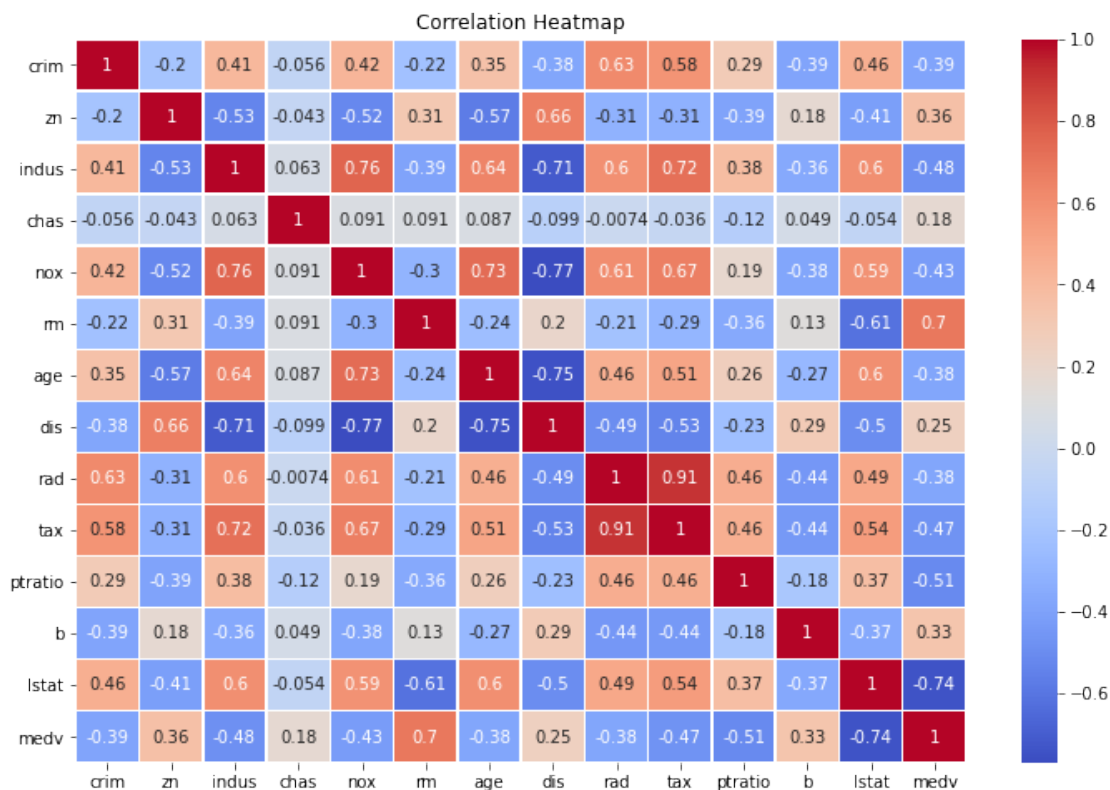

```

#	Column	Non-Null Count	Dtype
0	crim	506 non-null	float64
1	zn	506 non-null	float64
2	indus	506 non-null	float64
3	chas	506 non-null	int64
4	nox	506 non-null	float64
5	rm	506 non-null	float64
6	age	506 non-null	float64
7	dis	506 non-null	float64
8	rad	506 non-null	int64
9	tax	506 non-null	int64
10	ptratio	506 non-null	float64
11	b	506 non-null	float64
12	lstat	506 non-null	float64
13	medv	506 non-null	float64


dtypes: float64(11), int64(3)

memory usage: 55.5 KB

None



Mean Squared Error: 24.40482518814633



Predicted Price of the House: \$29.94

[]: